

*\*Experiment No. 4 \**

Aim: Design and implement a CNN for Image Classification

- Select a suitable image classification dataset (medical imaging, agricultural, etc.).
- Optimized with different hyper-parameters including learning rate, filter size, no. of layers, optimizers, dropouts, etc.

```
1 import tensorflow as tf
2 from tensorflow.keras.models import Sequential
3 from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
4 from tensorflow.keras.optimizers import Adam
5 from tensorflow.keras.datasets import cifar10
6 from tensorflow.keras.utils import to_categorical
7 import matplotlib.pyplot as plt
8 import seaborn as sns
9 from sklearn.metrics import confusion_matrix
10 import numpy as np
```

```
1 # Load CIFAR-10 dataset
2 (x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170498071/170498071 ————— 2s 0us/step

```
1 # Normalize pixel values
2 x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
1 # One-hot encode labels
2 y_train = to_categorical(y_train, 10)
3 y_test = to_categorical(y_test, 10)
```

```
1 # Build CNN model
2 model = Sequential([
3     Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
4     MaxPooling2D(2, 2),
5     Conv2D(64, (3, 3), activation='relu'),
6     MaxPooling2D(2, 2),
7     Conv2D(128, (3, 3), activation='relu'),
8     Flatten(),
9     Dense(128, activation='relu'),
10    Dropout(0.5),
11    Dense(10, activation='softmax')
12 ])
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape` to `input\_shape` in the super().\_\_init\_\_(activity\_regularizer=activity\_regularizer, \*\*kwargs)

```
1 # Compile model
2 model.compile(optimizer=Adam(learning_rate=0.001), loss='categorical_crossentropy', metrics=['accuracy'])
```

```
1 # Train model
2 epochs = 10
3 history = model.fit(x_train, y_train, epochs=epochs, validation_data=(x_test, y_test), batch_size=64)
```

Epoch 1/10  
782/782 ————— 76s 94ms/step - accuracy: 0.3030 - loss: 1.8829 - val\_accuracy: 0.5365 - val\_loss: 1.3018  
Epoch 2/10  
782/782 ————— 82s 95ms/step - accuracy: 0.5233 - loss: 1.3404 - val\_accuracy: 0.6167 - val\_loss: 1.0810  
Epoch 3/10  
782/782 ————— 82s 95ms/step - accuracy: 0.5973 - loss: 1.1492 - val\_accuracy: 0.6558 - val\_loss: 0.9809  
Epoch 4/10  
782/782 ————— 73s 94ms/step - accuracy: 0.6409 - loss: 1.0255 - val\_accuracy: 0.6798 - val\_loss: 0.9026  
Epoch 5/10  
782/782 ————— 83s 95ms/step - accuracy: 0.6749 - loss: 0.9338 - val\_accuracy: 0.6959 - val\_loss: 0.8750  
Epoch 6/10  
782/782 ————— 81s 93ms/step - accuracy: 0.6952 - loss: 0.8802 - val\_accuracy: 0.6786 - val\_loss: 0.9099  
Epoch 7/10  
782/782 ————— 81s 92ms/step - accuracy: 0.7187 - loss: 0.8119 - val\_accuracy: 0.6877 - val\_loss: 0.9233  
Epoch 8/10  
782/782 ————— 82s 92ms/step - accuracy: 0.7405 - loss: 0.7548 - val\_accuracy: 0.7228 - val\_loss: 0.8046  
Epoch 9/10  
782/782 ————— 82s 93ms/step - accuracy: 0.7511 - loss: 0.7190 - val\_accuracy: 0.7300 - val\_loss: 0.7867

Epoch 10/10

782/782 ————— 74s 95ms/step - accuracy: 0.7659 - loss: 0.6742 - val\_accuracy: 0.7364 - val\_loss: 0.7807

```

1 # Evaluate model
2 loss, accuracy = model.evaluate(x_test, y_test)
3 print(f"Test Accuracy: {accuracy * 100:.2f}%")

```

313/313 ————— 4s 14ms/step - accuracy: 0.7336 - loss: 0.7835  
 Test Accuracy: 73.64%

```

1 # Predictions
2 y_pred = model.predict(x_test)
3 y_pred_classes = np.argmax(y_pred, axis=1)
4 y_true_classes = np.argmax(y_test, axis=1)

```

313/313 ————— 4s 13ms/step

```

1 # Confusion matrix
2 cm = confusion_matrix(y_true_classes, y_pred_classes)
3 plt.figure(figsize=(8, 6))
4 sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
5 plt.xlabel('Predicted Label')
6 plt.ylabel('True Label')
7 plt.title('Confusion Matrix')
8 plt.show()

```



